

Kubernetes Networking 101

Shubham Sharma
Microsoft

@shubham1172



Agenda

Container-to-Container communication

Pod-to-Pod communication

Service discovery

Looking beyond the cluster

What next?





Agenda

Container-to-Container communication

Pod-to-Pod communication

Service discovery

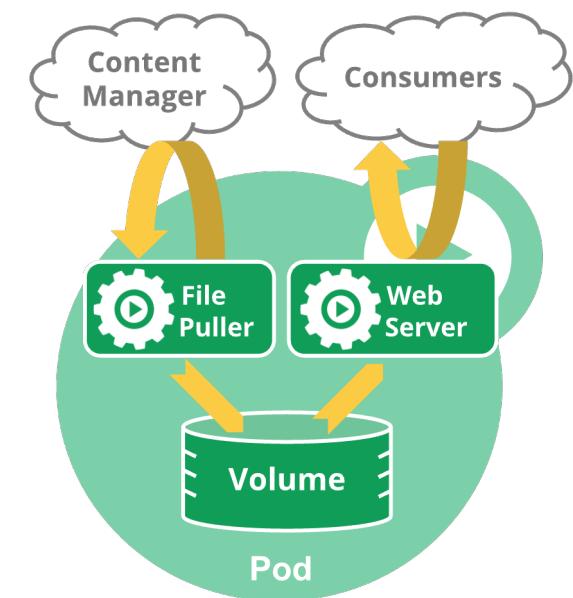
Looking beyond the cluster

What next?



Pods

- > Smallest deployable units of computing in Kubernetes
- > Collection of one or more containers
- > Unique IP address
- > **Shared networking and storage**



[\[1\] How Pods manage multiple containers](#)

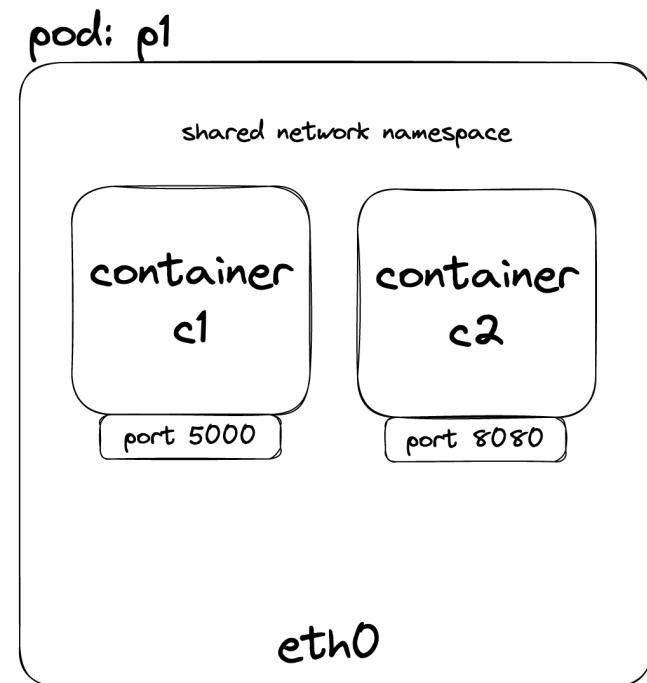


Container networking

Every container in a Pod shares the network namespace, including the IP address and network ports.

Containers can communicate over

- localhost
- IPC (like POSIX shared memory)





Agenda

Container-to-Container communication

Pod-to-Pod communication

Service discovery

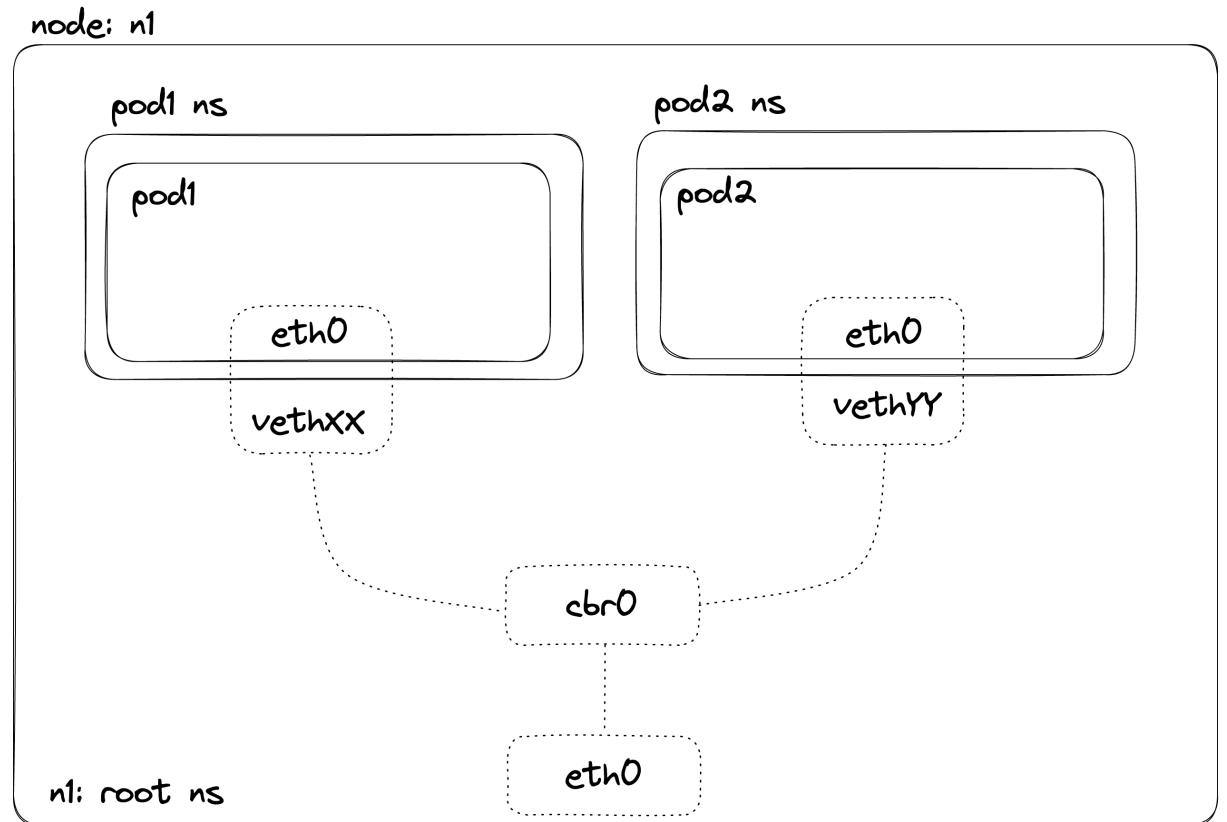
Looking beyond the cluster

What next?



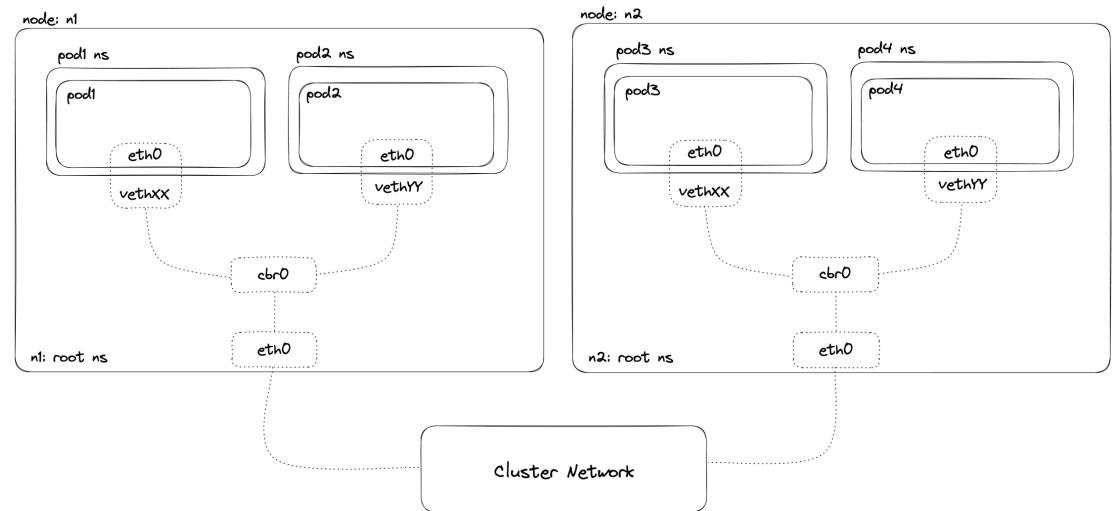
Pods on the same node

- > Each pod has an IP address
- > Virtual Ethernet Pairs
- > Linux Ethernet Bridge (cbr0, docker0)



Pods across multiple nodes

- > Bridge fails to resolve the MAC address
- > Container Network Interface (CNI)
- > CNI allocates IPs to Pods
- > Different implementations (layer 2, layer 3, overlay, vxlan)





Agenda

Container-to-Container communication

Pod-to-Pod communication

Service discovery

Looking beyond the cluster

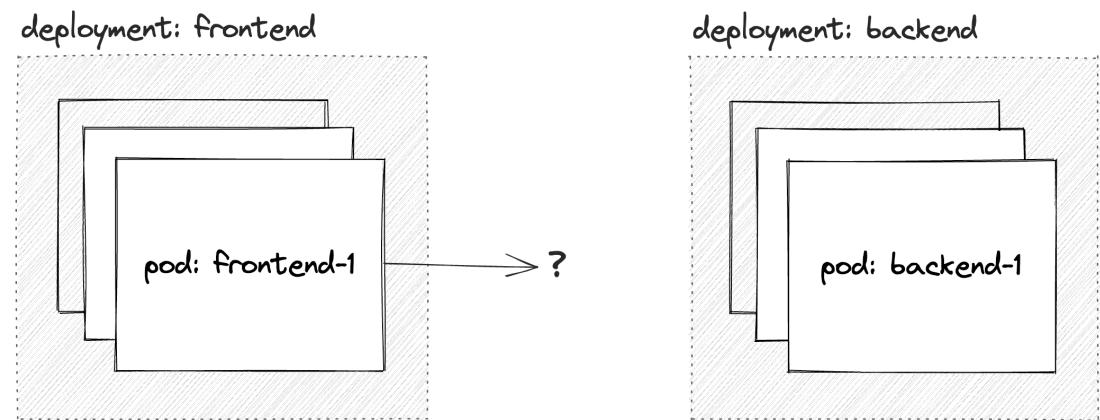
What next?



Service Discovery 101

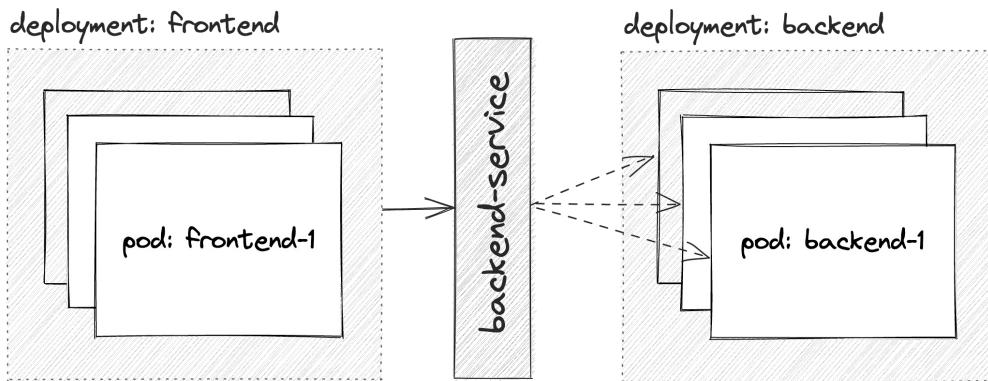
- > Virtual IPs are dynamic.
- > Pods can be created/destroyed dynamically
- > Which Pods are healthy?

How can a frontend Pods keep track of which IP addresses to connect to?



Service Object

- > Service “backend-service” targets TCP **port 8080** on any Pod **with label** “app.kubernetes.io/name: backend”
- > Service has an IP address (**ClusterIP**)
- > Provides load-balancing across Pods



```
apiVersion: v1
kind: Pod
metadata:
  name: backend
  labels:
    app.kubernetes.io/name: backend
spec:
  containers:
  - name: mybackend
    image: contoso/mybackend
    ports:
    - containerPort: 8080
---
apiVersion: v1
kind: Service
metadata:
  name: backend-service
spec:
  selector:
    app.kubernetes.io/name: backend
  ports:
  - protocol: TCP
    port: 80
    targetPort: 8080
```



Service DNS

- > Cluster IP is virtual and stable
- > Kubernetes DNS service
- > Search domains

```
# DNS A record of K8s service  
service.namespace.svc.cluster.local  
  
# Contacting a service in the same namespace  
other-service  
  
# Contacting a service in another namespace  
other-service.other-namespace
```



Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: backend
spec:
  replicas: 3
  selector:
    matchLabels:
      app.kubernetes.io/name: backend
  template:
    metadata:
      labels:
        app.kubernetes.io/name: backend
    spec:
      containers:
        - name: mybackend
          image: contoso/mybackend
          ports:
            - containerPort: 8080
```

```
apiVersion: v1
kind: Service
metadata:
  name: backend-service
spec:
  selector:
    app.kubernetes.io/name: backend
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: frontend
spec:
  replicas: 3
  selector:
    matchLabels:
      app.kubernetes.io/name: frontend
  template:
    metadata:
      labels:
        app.kubernetes.io/name: frontend
    spec:
      containers:
        - name: myfrontend
          image: contoso/myfrontend
          env:
            - name: BACKEND_SERVICE_NAME
              value: backend-service
```

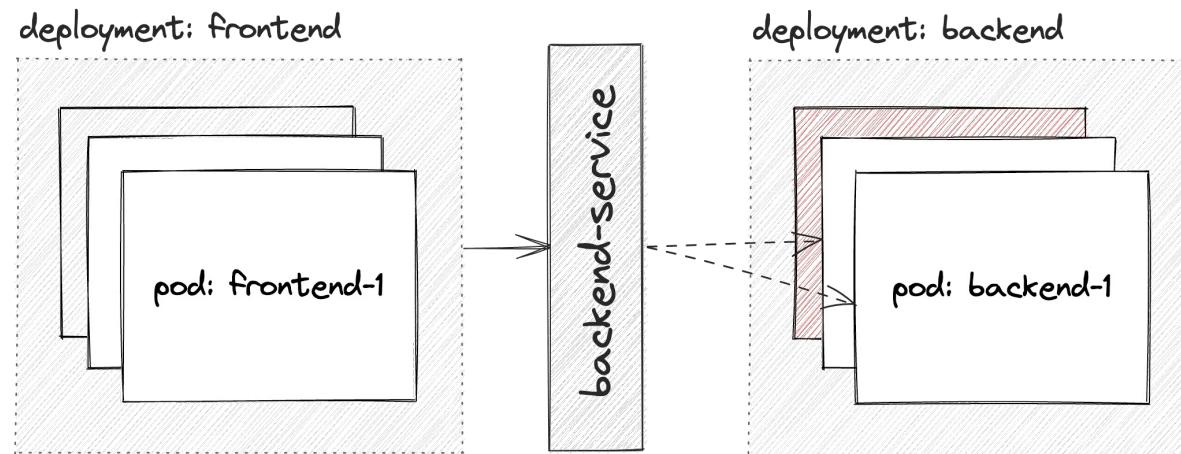


Readiness checks

Service objects can track which Pods are ready via a readiness check.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: backend
spec:
  ...
  template:
    ...
    spec:
      containers:
        - name: mybackend
          image: consoto/mybackend
          readinessProbe:
            httpGet:
              path: /health
              port: 8080
            initialDelaySeconds: 5
            periodSeconds: 10
            failureThreshold: 3
  ...

```





Agenda

Container-to-Container communication

Pod-to-Pod communication

Service discovery

Looking beyond the cluster

What next?



Incoming traffic?

- > ClusterIPs are virtual IPs **within** the cluster
- > How to allow incoming traffic to Pods?

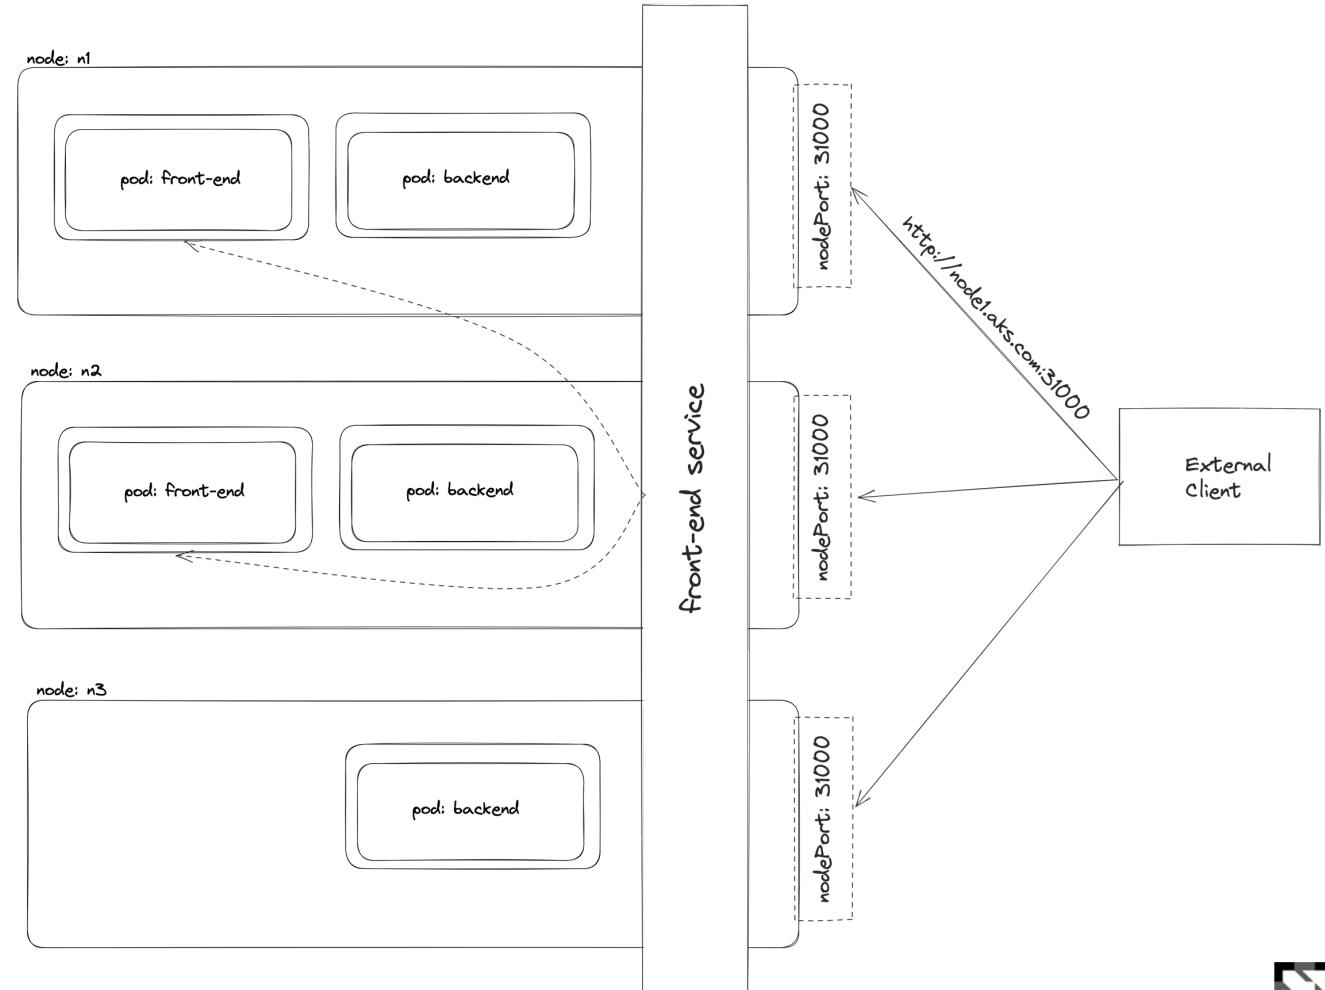


NodePort

Allow access from external clients via a universal port –

NodePort

- > Builds on top of Cluster IP internally
- > Every node listens on a port and proxies it into a Service
- > Port is reserved cluster-wide



NodePort

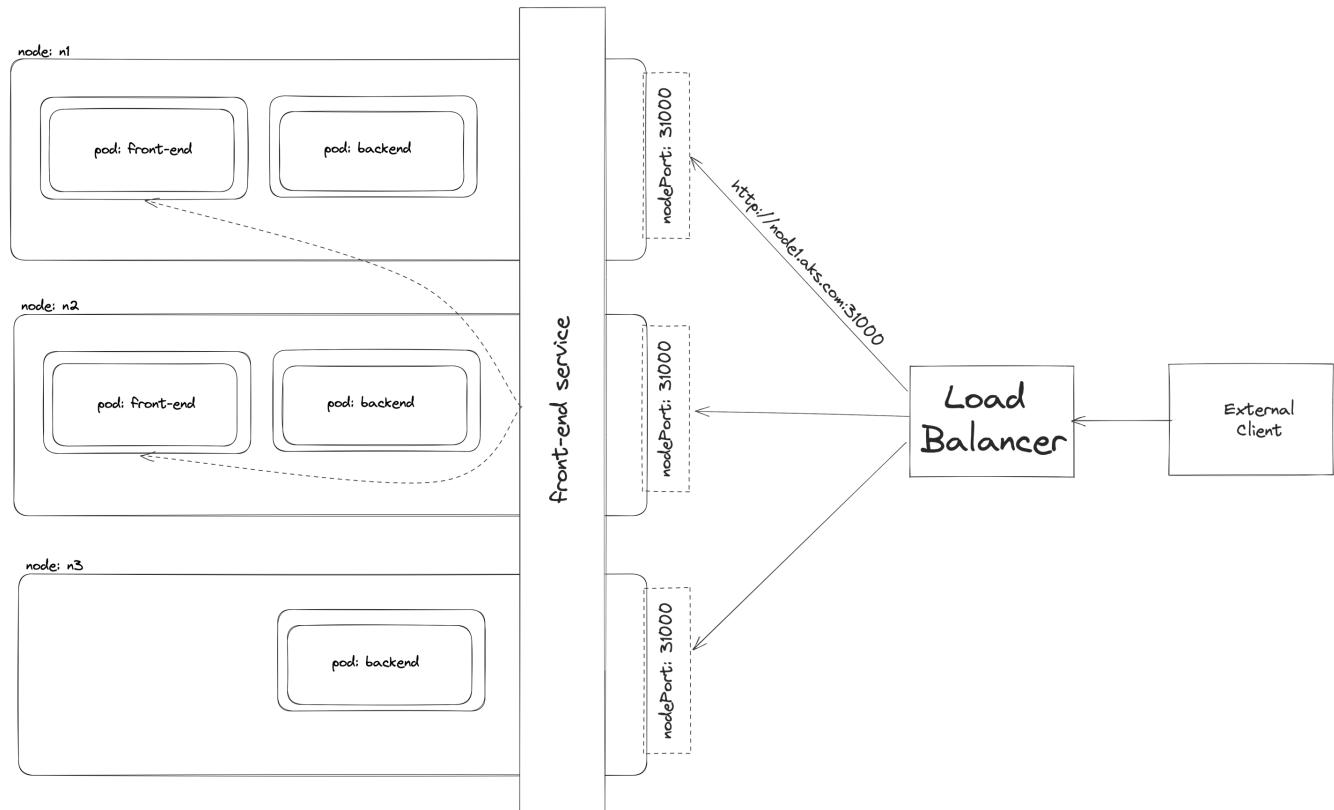
```
apiVersion: v1
kind: Service
metadata:
  name: frontend-service
spec:
  type: NodePort
  selector:
    app.kubernetes.io/name: frontend
  ports:
    - port: 80
      targetPort: 8080
      # nodePort is exposed on all nodes in the cluster
      nodePort: 31000
```



LoadBalancer

Provisions an **external** load balancer for your service (based on the cloud provider)

- > Builds on top of NodePort to expose pods internally
- > Administrators can configure scaling, firewall, routing, etc. on the external load balancer



LoadBalancer

```
apiVersion: v1
kind: Service
metadata:
  name: frontend-service
spec:
  selector:
    app.kubernetes.io/name: frontend
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
  type: LoadBalancer
```



Services without selector

How to access services that run outside the cluster?

> Use service discovery via DNS for external resources

> External database in production, internal in testing

```
apiVersion: v1
kind: Service
metadata:
  name: external-database
spec:
  # Remember, no selectors!
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

```
apiVersion: discovery.k8s.io/v1
kind: EndpointSlice
metadata:
  name: external-database-1 # by convention, use the name of the Service
                                # as a prefix for the name of the EndpointSlice
  labels:
    # You should set the "kubernetes.io/service-name" label.
    # Set its value to match the name of the Service
    kubernetes.io/service-name: external-database
spec:
  addressType: IPv4
  ports:
    - name: ''
      appProtocol: http
      protocol: TCP
      port: 9376
  endpoints:
    - addresses:
        - "10.4.5.6" # the IP addresses in this list can appear in any order
        - "10.1.2.3"
```





Agenda

Container-to-Container communication

Pod-to-Pod communication

Service discovery

Looking beyond the cluster

What next?

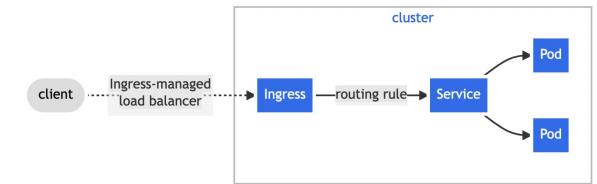


HTTP load balancing with Ingress

Service object operates at L4,
how to route based on HTTP
(L7)?

Ingress: Kubernetes native way
to do “virtual hosting” – host
many sites on a single IP.

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-ingress
spec:
  rules:
    - host: "foo.bar.com"
      http:
        paths:
          - pathType: Prefix
            path: "/bar"
            backend:
              service:
                name: service1
                port:
                  number: 80
    - host: "*.foo.com"
      http:
        paths:
          - pathType: Prefix
            path: "/foo"
            backend:
              service:
                name: service2
                port:
                  number: 80
```



[2] [What is ingress?](#)

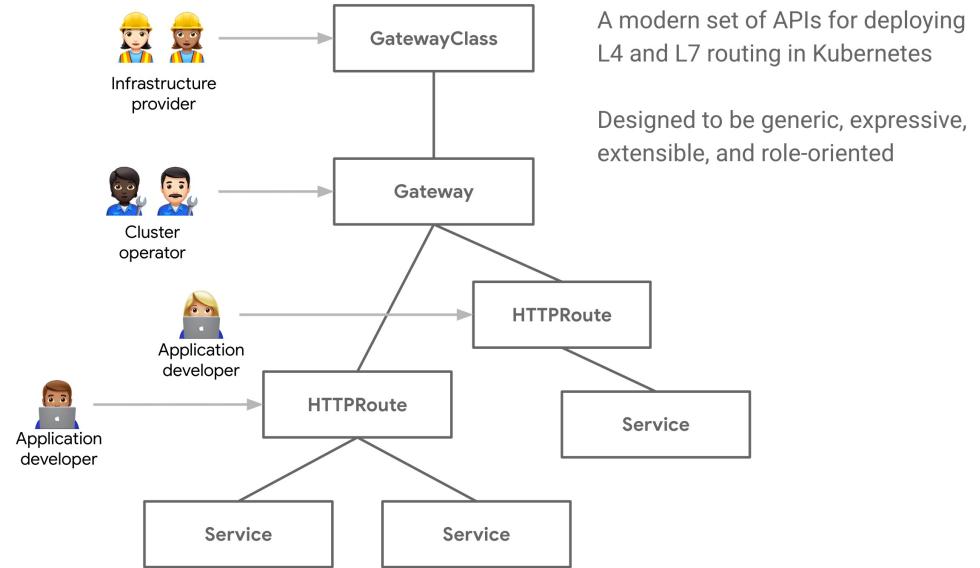


Gateway API

While Ingress is a useful abstraction, it has not scaled very well.

- > Some features are underdefined, implementations may vary, reducing portability.
- > Easy to misconfigure it.

Gateway API (being developed by SIG-network) is working on a modern API for L7 (and even L4) routing.



[\[3\]: What is the Gateway API?](#)



More topics

Kube-proxy

CNI types and implementations

Endpoints and EndpointSlices

Ingress Controllers

Service Meshes



References

- [1] <https://kubernetes.io/docs/concepts/workloads/pods/#pod-networking>
- [2] <https://kubernetes.io/docs/concepts/services-networking/ingress/#what-is-ingress>
- [3] <https://gateway-api.sigs.k8s.io/#what-is-the-gateway-api>



Thank you!